

CSCI 5253 Final Project Report

Zane Jakobs

December 9, 2020

1 Overview And Goals

The main purpose of this project is to provide an extensible framework for the use of Extended Berkeley Packet Filter (eBPF)-based network profiling in MPI clusters. If a programmer goes to the effort of writing and debugging an MPI program—or even if the user simply wants to compare among existing programs—they will often care about that program running fast and efficiently on a cluster. Often, network bandwidth is the biggest bottleneck of such programs, particularly in scientific computing, [INSERT CITE HERE] so a programmer interested in a high-performance application will want to be able to monitor their application’s network usage, ideally with minimal CPU and network overhead. Libraries like PETSc enable such tracking easily — for instance, any PETSc program can be run with the command-line option `–log-view` to view a summary of the program’s performance statistics, including network usage. However, MPI programmers that are not using PETSc or a similar library generally will have to recompile their application to get a network usage profile or trace from tools like VampirTrace or PMPI. As anyone who has had to compile a large scientific application knows, this is not always as straightforward as it sounds, and can take a long time even when the process goes smoothly.

One solution to this problem comes from the modern linux kernel, which enables very low-level tracing of both kernel- and user-space events by attaching eBPF programs to kprobes. [CITE Ingo Molnar <https://lkm1.org/lkm1/2015/4/14/232>] The BPF Compiler Collection (BCC) [CITE BCC Github] provides a set of pre-written eBPF programs that are useful for instrumentation of all sorts, including network profiling. In particular, the programs `tcpaccept(8)`, `tcpconnect(8)`, `tcpconnlst(8)`, `tcplife(8)`, and `tcpretrans(8)` can be combined to produce a useful summary of a program’s network usage.

Specifically, if each node in an MPI cluster is running all of the above-named eBPF programs (`tcpaccept(8)` and `friends`), their outputs can be combined together to produce process-wise summaries of TCP traffic, which we store in a hash table whose keys are process ID numbers (PIDs). Periodically, each node pushes its summaries into an MPI-backed message queue, and those queues are sent back to the root node. After receiving those messages containing summaries on other MPI ranks (each running on its own node), the root node writes them to an output file. Additionally, as soon as data is available, the root node launches a Flask server that serves requests to a REST API that allows the user to request all collected data, all data on a certain MPI rank with a given PID, or all data for all processes (running on any MPI rank) with a given name.

2 Description of Components

The project can be broadly subdivided into 7 components: the eBPF programs, data structures to serialize and efficiently encode their output and summaries of it, an MPI-backed message queue, a REST API and server, the driver programs, an optional containerized version of the program (hosted on Singularity hub) and for testing purposes, cluster setup and installation scripts to build an SDN, specifically a LAN. Since the project is, with the exception of the REST server, written in C, and the PETSc library provides many of the utilities programmers expect from “modern” languages that C doesn’t provide, as well as an easy interface to install many useful C libraries through its configure script (and update the relevant environment variables to use those libraries in a makefile), we make heavy use of PETSc in our code. For example, we use a hash map in the form of a PETSc wrapper around a KHash hash map, and our message queue stores `PetscBag` instances, which can in turn store any of the data and summary structs we define.

eBPF Programs

Due to time constraints, we used only the five BCC-provided programs mentioned in the first section: `tcpaccept(8)`, `tcpconnect(8)`, `tcpconnlst(8)`, `tcplife(8)`, and `tcpretrans(8)`. Through a script named `collect -tcp -data.sh` that must, at present, be run as root¹, each of these programs is run in the

¹Since all eBPF programs must be run as root, although this may change in the future. [CITE Ingo Molnar <https://lkml.org/lkml/2015/4/14/232>]

background with its output redirected to an appropriately named file. This must be done on every node in the cluster before running the program to be profiled; this can be accomplished with OpenMPI on N nodes by running

```
sudo mpirun -np N --map-by ppr:1:node --allow-run-as-root ./collect-tcp-data.sh &
```

from the root node. Every MPI implementation I’ve seen will warn you that you should not run as root or outright disallows it by default. In OpenMPI, this can be circumvented with the `--allow-run-as-root` option. The requirement that the data collection script be run as root presents the most significant limitation of this software: it must be run on a cluster to which the user has root access.

We also define a C struct for each program that can hold the information in one line of output, as well as a C struct to hold data from all the programs collected about an individual process, and one to hold a summary of that process’s data.

Serialization

To allow the user to extend the program to collect more sophisticated summaries, each of the above-mentioned C structs can be sent in an MPI message after the user has called the `register_mpi_types()` function. This was accomplished with MPI’s `MPI_Type_create_struct()` and `MPI_Type_commit()` functions. By default, the output summaries are stored in plain ASCII text in a file. However, if the user has extremely limited disk space and wants to use a more efficient encoding, then as mentioned above, each of those structs can be serialized in a `PetscBag`, which is a C class provided by PETSc that can be used for, among other things, efficiently serializing data in a binary format. Additionally, for testing purposes we wrote a Python script to download matrices from the Matrix Market (math.nist.gov/MatrixMarket) and serialize them into the space-efficient PETSc binary format.

The Message Queue

Our extremely simple message queue is built based on a circular buffer and the `MPI_Reduce()` and `MPI_Gather()` functions. We decided that it would be easier, both for development but crucially for the end user, to implement a simple message queue with MPI than to integrate a “standard” message queue like `RabbitMQ` or `ZeroMQ` into our cluster and program setup. Its operation consists of all nodes (possibly including the root node) collecting data and pushing it into their respective buffers, then gathering that data

to the root node, clearing the non-root nodes' buffers. Since this operation requires the use of `MPI_Reduce()` (to figure out how many items each process is sending so that the root node can allocate enough memory to store them) and `MPI_Gather()` (to perform the actual operation), both of which are blocking², the user does not need to worry about using `MPI_Barrier()` to synchronize the communication. This is implemented in the `buffer_gather()` function in our code.

Driver Programs

There are two C programs that run the code contained in `petsc_webserver.h`, `petsc_webserver.c`, and `petsc_webserver_backend.py`: `petsc_webserver_driver.c` and `webserver_launcher.c`. The main driver first reads through all the relevant input files and stores the data it reads. It then uses `MPI_Comm_spawn()` to spawn a subprocess on the root node (running `webserver_launcher.c`) that launches the REST server via the Python C API (specifically, with `PyRun_SimpleFile()`). The driver then continues into an infinite loop, where it polls the data files to check if they have new data, and if they do, reads that new data, updates the summaries, then overwrites the old output file with the updated output. The driver program can be run with the following options, produced by running it with the `-h` (or `-help`) flag (excluding options common to all PETSc programs, which are also printed out):

PETSc webserver: This program periodically reads the output of the eBPF programs `tcpaccept`, `tcpconnect`, `tcpconnlat`, `tcplife`, and `tcpretrans`, summarizes that data, and stores that data in a message queue. The summaries are then gathered to the root node (MPI rank 0), where it is written to an output file. Once the existing input files have been read to their ends, the root process spawns a new subcommunicator via `MPI_Comm_spawn()` and launches a Python webserver (currently written with Flask) on that spawned process that reads the output file, and serves requests to a REST API. Available endpoints are:

```
-----
GET /api/get/all (gets data for all processes on all MPI ranks)
GET /api/get/{rank}/{pid} (gets data for the process with PID {pid} on MPI rank {rank})
GET /api/get/{name} (gets data on all MPI ranks for all processes with name {name})
-----
```

Usage:

Options:

`--python_server [filename]` : (optional, default `petsc_webserver_frontend.py`) filename of Python web

²i.e. they block all processes in the communicator until every other process has finished executing the function

you want to launch

```
--python_launcher [filename] : (optional, default ./webserver_launcher) filename of the MPI program
    that can be launched as an MPI child with MPI_Comm_split() with argv
    'python3 <python_server> -f <output> -p <port>' and will run a webserver on the
    appropriate port.
-file [filename] : (optional if any of --XXX_file are given) input file
-type [TCPACCEPT,TCPCONNECT,TCPRETRANS,TCPLIFE,TCPCONNLAT] : (required only if -file is given) what
    sort of input file is the -file argument?
--accept_file [filename] : (optional) file for tcpaccept data
--connect_file [filename] : (optional) file for tcpconnect data
--connlat_file [filename] : (optional) file for tcpconnlat data
--life_file [filename] : (optional) file for tcplife data
--retrans_file [filename] : (optional) file for tcpretrans data
-o (--output) [filename] : (optional, default stdout) file to output data to
-p (--port) [port (int)] : (optional, default 5000) which TCP port to use to serve requests
--buffer_capacity [capacity] : (optional, default 10,000) size of the buffer (number of entries)
--polling_interval [interval] : (optional, default 5.0) how many seconds to wait before
    checking the file for more data after reaching the end?
```

The root endpoint displays HTML representing all the data we have so far, and the endpoint /<name> retrieves HTML representing all entries for the given program name. Note that the driver MUST be started after TCP data collection has started in order to guarantee that every file the driver reads in exists.

I chose this design because I want this software to be as portable as it can be across different clusters, so I tried to rely as little as possible on anything outside of the C standard, POSIX (since eBPF requires a modern linux kernel anyway), the MPI standard or PETSc. However, one notable drawback is that I can't do online updates of the server frontend because it's launched from MPI_Comm_spawn within the backend, so updating the frontend requires restarting the backend (which automatically launches the frontend). It would be possible to make this happen automatically by providing a git hook to install on the server so that on each commit, git rebuilds the server if the backend changes and restarts it if either the backend or frontend changes, but we did not do this.

REST API

The Python program launched in the drivers above takes in two command-line options, one (-f) being the file containing its input data (the output of the driver), the other (-p) being the TCP port that Flask should use.

The endpoints for the REST API are shown in the help message above. Upon receiving a request, the server re-reads the input data (in case it has been updated since the last request), constructs the appropriate response, and responds. I originally intended to use SAWs (Scientific Application Webserver), a library that PETSc integrates with that would have allowed me to bypass the output file and subprocess running the server entirely. However, the PETSc developers haven't written the code necessary to use PetscBag objects with SAWs (which is what I was planning to do), so for the sake of time, since implementing even a very simple function for a library like PETSc can be a significant time investment, we decided to use a Flask-based server for now. In case the user doesn't feel like writing their own REST calls (even though this API is extremely simple, and requests can easily be constructed manually), we have provided a Python-based web client.

Singularity Container

In case the user doesn't want to build the software on their cluster, we have provided the software in the form of a Singularity container at [INSERT SHUB LINK]. Also, if it ever becomes the case that eBPF programs can be run without root privileges, this container will allow the software to be run on clusters whose users do not have root access. For ease of use on systems that use Modules, we also provide a Tool Command Language (TCL) modulefile to provide access to Singularity.

Cluster Setup Scripts and SD-LAN

In addition to providing a containerized version of the software, we wrote some shell scripts to install the relevant dependencies and software on an Ubuntu VM. To test the software, we set up two VMs in Chameleon Cloud running Ubuntu 20.04, and, following the guide at <https://mpitutorial.com/tutorials/running-an-mpi-cluster-within-a-lan/>, set up a LAN between the two with ssh-agent, with a distributed filesystem running NFS that we used to retrieve debugging information from the non-root node, as well as distributing the executable among the nodes.

3 Capabilities and Limitations

This project provides a low-overhead³ option to profile the network usage of an MPI program and transmit the results to a remote user via a REST API (including browser-friendly HTML). It does not require the user to recompile their code, only to run the data collection programs and then the server before running their program. This could be done explicitly on the command-line or in a batch script, running the data collection programs and server in the background. Furthermore, even an inexperienced user should be able to extend this code to collect more detailed data if so desired.

Appendix A

This code — available at <https://github.com/DiffeoInvariant/datacenter-profiling> — defines the public API of the mini-library that powers the driver program, itself available along with the implementation of this header and the Python webserver code in that git repository.

```
1 #ifndef DCPROF_PETSC_WEBSERVER_H
2 #define DCPROF_PETSC_WEBSERVER_H
3 #include <petsc.h>
4 #include <unistd.h>
5 #include <petsc/private/hashtable.h>
6 #include <petsc/private/hashmap.h>
7
8 #define IP_ADDR_MAX_LEN 45
9 #define COMM_MAX_LEN PETSC_MAX_PATH_LEN
10
11 typedef enum {TCPACCEPT, TCPCONNECT, TCPCONNLAT, TCPLIFE, TCPRETRANS} InputType;
12 static const char *InputTypes[] = {"ACCEPT", "CONNECT", "CONNLAT", "LIFE", "RETRANS", "TCP", 0};
13
14 typedef enum
15 {
16     DTYPE_ACCEPT=0,
17     DTYPE_CONNECT=1,
18     DTYPE_CONNLAT=2,
19     DTYPE_LIFE=3,
20     DTYPE_RETRANS=4,
21     DTYPE_SUMMARY=5
22 } SERVER_MPI_DTYPE;
23
24 MPI_Datatype MPI_DTYPES[6];
```

³Even if the data collecting programs have generated over 20,000 total entries, the program uses less than 500 kB of network data (i.e. transmits \leq 500 kB over the network).

```

25
26 /* call this at the beginning of the program, but after MPI_Init(), for
27    any program that needs to send any of the XXX_entry types here, as well
28    as process_data_summary. The registered types are stored in MPI_DTYPES
29    and indexed by the SERVER_MPI_DTYPES enum. */
30 extern PetscErrorCode register_mpi_types();
31
32 typedef struct {
33     PetscInt pid, ip, rport, lport;
34     char      laddr[IP_ADDR_MAX_LEN], raddr[IP_ADDR_MAX_LEN], comm[COMM_MAX_LEN];
35 } tcpaccept_entry;
36
37 /* creates a PetscBag with PetscBagCreate() to serialize a
38    tcpaccept_entry, and stores a pointer to that entry
39    (via PetscBagGetData()) in the first parameter. The second parameter
40    is used to store the PetscBag we create. The third parameter indicates
41    how many entries you have created already to ensure that each entry has
42    a unique name, even if data like the PID and/or process name are the
43    same. */
44 extern PetscErrorCode create_tcpaccept_entry_bag(tcpaccept_entry **, PetscBag *,
45     PetscInt);
46
47 /* parses a line of the form
48    PID COMM IP RADDR RPORT LADDR LPORT
49    which is stored in the second parameter,
50    and the result is written to the first parameter */
51 extern PetscErrorCode tcpaccept_entry_parse_line(tcpaccept_entry *, char *);
52
53 typedef struct {
54     PetscInt pid, ip, dport;
55     char      saddr[IP_ADDR_MAX_LEN], daddr[IP_ADDR_MAX_LEN], comm[COMM_MAX_LEN];
56 } tcpconnect_entry;
57
58 /* creates a PetscBag with PetscBagCreate() to serialize a
59    tcpconnect_entry, and stores a pointer to that entry
60    (via PetscBagGetData()) in the first parameter. The second parameter
61    is used to store the PetscBag we create. The third parameter indicates
62    how many entries you have created already to ensure that each entry has
63    a unique name, even if data like the PID and/or process name are the
64    same. */
65 extern PetscErrorCode create_tcpconnect_entry_bag(tcpconnect_entry **, PetscBag *,
66     PetscInt);
67
68 /* parses a line of the form
69    PID COMM IP SADDR DADDR DPORT
70    which is stored in the second parameter,
71    and the result is written to the first parameter */
72 extern PetscErrorCode tcpconnect_entry_parse_line(tcpconnect_entry *, char *);
73

```



```

72 typedef struct {
73     PetscInt  pid,ip,dport;
74     PetscReal lat_ms;
75     char      saddr[IP_ADDR_MAX_LEN],daddr[IP_ADDR_MAX_LEN],comm[COMM_MAX_LEN];
76 } tcpconnlat_entry;
77
78 /* creates a PetscBag with PetscBagCreate() to serialize a
79 tcpconnlat_entry, and stores a pointer to that entry
80 (via PetscBagGetData()) in the first parameter. The second parameter
81 is used to store the PetscBag we create. The third parameter indicates
82 how many entries you have created already to ensure that each entry has
83 a unique name, even if data like the PID and/or process name are the
84 same. */
85 extern PetscErrorCode create_tcpconnlat_entry_bag(tcpconnlat_entry **, PetscBag *,
86     PetscInt);
87
88 /* parses a line of the form
89 PID COMM IP SADDR DADDR DPORT LAT(ms)
90 which is stored in the second parameter,
91 and the result is written to the first parameter */
92 extern PetscErrorCode tcpconnlat_entry_parse_line(tcpconnlat_entry *, char *);
93 #define TIME_LEN 9
94
95 typedef struct {
96     PetscInt  pid,ip,lport,rport,tx_kb,rx_kb;
97     PetscReal ms;
98     char      laddr[IP_ADDR_MAX_LEN],raddr[IP_ADDR_MAX_LEN],comm[COMM_MAX_LEN],time[
99         TIME_LEN];
100 } tcplife_entry;
101
102 /* creates a PetscBag with PetscBagCreate() to serialize a
103 tcplife_entry, and stores a pointer to that entry
104 (via PetscBagGetData()) in the first parameter. The second parameter
105 is used to store the PetscBag we create. The third parameter indicates
106 how many entries you have created already to ensure that each entry has
107 a unique name, even if data like the PID and/or process name are the
108 same. */
109 extern PetscErrorCode create_tcplife_entry_bag(tcplife_entry **, PetscBag *,
110     PetscInt);
111
112 /* parses a line of the form
113 PID,COMM,IP,LADDR,LPORT,RADDR,RPORT,TX_KB,RX_KB,MS
114 which is stored in the second parameter,
115 and the result is written to the first parameter.
116 NOTE: entries of this form can be generated by
117 running 'tcplife -s > /path/to/tcplife.data.file',
118 since they are comma-delimited, unlike the other
119 XXX_entry_parse_line() functions declared in this file
120 whose input is space-delimited. */

```

```

118 extern PetscErrorCode tcplife_entry_parse_line(tcplife_entry *, char *);
119
120
121 typedef struct {
122     PetscInt pid,ip;
123     char      laddr_port[COMM_MAX_LEN],raddr_port[COMM_MAX_LEN],state[COMM_MAX_LEN];/*
        TODO: find out how long these really should be, cuz this is longer than
        necessary. not too important though. */
124 } tcpretrans_entry;
125
126 /* creates a PetscBag with PetscBagCreate() to serialize a
127    tcpretrans_entry, and stores a pointer to that entry
128    (via PetscBagGetData()) in the first parameter. The second parameter
129    is used to store the PetscBag we create. The third parameter indicates
130    how many entries you have created already to ensure that each entry has
131    a unique name, even if data like the PID and/or process name are the
132    same. */
133 extern PetscErrorCode create_tcpretrans_entry_bag(tcpretrans_entry **, PetscBag *,
        PetscInt);
134
135
136 extern PetscErrorCode tcpretrans_entry_parse_line(tcpretrans_entry *, char *);
137
138
139 typedef struct {
140     /* a circular buffer */
141     PetscBag *buf; /* array of items */
142     size_t capacity,valid_start,valid_end,num_items;
143 } entry_buffer;
144
145 /* creates a circular buffer that can be used as a message queue,
146    with MPI for message passing. All MPI ranks can read data and store
147    it in the buffer. While you can store any type in the buffer that can
148    be serialized in a PetscBag, and on any one rank you can in principle
149    store multiple types in the same buffer, you CANNOT store multiple types
150    in the same buffer and then call buffer_gather() or buffer_gather_summaries()
151    on that buffer, as that invokes undefined behavior in MPI_Gather()
152
153    The first parameter is a pointer to the created buffer, and the second
154    parameter is the capacity of the buffer.*/
155 extern PetscErrorCode buffer_create(entry_buffer *, size_t);
156
157 /* frees all memory used by the buffer, calls PetscBagDestroy() on any
158    remaining elements */
159 extern PetscErrorCode buffer_destroy(entry_buffer *);
160
161 /* inserts the PetscBag stored in the second parameter in the buffer
162    pointed to by the first. Returns 0 on success, non-zero on failure.
163    The only possible failure mode is if the buffer is full, in which case

```

```

164     the insertion does not happen. */
165 extern PetscInt      buffer_try_insert(entry_buffer *, PetscBag);
166
167 /* returns the current number of items in the buffer pointed to by the first
   parameter */
168 extern size_t        buffer_size(entry_buffer *);
169
170 extern size_t        buffer_capacity(entry_buffer *);
171
172 extern PetscBool     buffer_full(entry_buffer *);
173
174 extern PetscBool     buffer_empty(entry_buffer *);
175
176 /* gets the object that has spent the most time in the buffer.
   If you call buffer_get_item() multiple times in a row WITHOUT
   calling buffer_pop() in between each call, you will get the same
   PetscBag on each call to buffer_get_item()*/
177
178 /* remove the object that has spent the most time in the buffer,
   and free its memory */
179
180 extern PetscErrorCode buffer_get_item(entry_buffer *, PetscBag *);
181
182 extern PetscErrorCode buffer_pop(entry_buffer *);
183
184 /* gathers all buffer summaries to a buffer on root */
185
186 extern PetscErrorCode buffer_gather_summaries(entry_buffer *);
187
188 /* gather everything in the buffer pointed to by the first parameter to root,
   with all the entries of the type indicated by the second parameter.
   NOTE: if all the entries are not of that type, that will invoke
   undefined behavior in the MPI_Gather(), and probably cause a crash */
189
190 extern PetscErrorCode buffer_gather(entry_buffer *, SERVER_MPI_DTYPE);
191
192
193
194
195
196 typedef struct {
197     long long naccept, nconnect, nconnlst, nlife, nretrans,
198     tx_kb, rx_kb, nipv4, nipv6;
199     PetscReal latms, lifems;
200     char comm[COMM_MAX_LEN];
201 } process_data;
202
203 typedef struct {
204     PetscInt pid, rank;
205     long tx_kb, rx_kb, n_event;
206     PetscReal avg_latency, avg_lifetime, fraction_ipv6;
207     char comm[COMM_MAX_LEN];
208 } process_data_summary;
209
210 /* write the summary pointed to by the second parameter to the
   file pointed to by the first */

```

```

212 extern PetscErrorCode summary_view(FILE *, process_data_summary *);
213
214 /* create a process_data_summary and store it in the third parameter. The
215    first parameter is the PID, and the second is a pointer to the
216    process_data you want to summarize.*/
217 extern PetscErrorCode process_data_summarize(PetscInt, process_data *,
        process_data_summary *);
218
219 #define pid_hash(pid) pid
220
221
222 #define process_data_equal(lhs,rhs) ( lhs.naccept == rhs.naccept && \
223     lhs.nconnect == rhs.nconnect && \
224     lhs.nconnlat == rhs.nconnlat && \
225     lhs.nlife == rhs.nlife && \
226     lhs.tx_kb == rhs.tx_kb && \
227     lhs.rx_kb == rhs.rx_kb && \
228     lhs.nipv4 == rhs.nipv4 && \
229     lhs.nipv6 == rhs.nipv6 && \
230     lhs.latms == rhs.latms && \
231     lhs.lifems == rhs.lifems)
232
233 #define int_equal(lhs,rhs) (lhs == rhs)
234
235 static process_data default_pdata = {0,0,0,0,0,0,0,0,0,0.0,0.0,"[unknown]"};
236
237 PETSC_HASH_MAP(HMapData,PetscInt,process_data,PetscHashInt,int_equal,default_pdata)
        ;
238
239 /* sets the values of the process_data to the defaults */
240 extern PetscErrorCode process_data_initialize(process_data *);
241
242 /* calculates what fraction of all TCP events used IPv6 as opposed to IPv4 */
243 extern PetscReal fraction_ipv6(process_data *);
244
245 extern PetscErrorCode create_process_data_bag(process_data **, PetscBag *);
246
247 /* the third parameter is the MPI_Comm rank, fourth is the number of the entry*/
248 extern PetscErrorCode create_process_summary_bag(process_data_summary **, PetscBag
        *, PetscInt, PetscInt);
249
250 typedef struct {
251     PetscHMapData ht;
252 } process_statistics;
253
254 extern PetscErrorCode process_statistics_get_summary(process_statistics *, PetscInt
        , process_data_summary *);
255

```

```

256 extern PetscErrorCode process_statistics_num_entries(process_statistics *, PetscInt
    *);
257
258 /* second and third parameters are pointers to array of process_data and PetscInt (
    pid) respectively, each of the length retrieved from
    process_statistics_num_entries() */
259 extern PetscErrorCode process_statistics_get_all(process_statistics *, process_data
    *, PetscInt *);
260
261
262 extern PetscErrorCode process_statistics_init(process_statistics *);
263
264 extern PetscErrorCode process_statistics_destroy(process_statistics *);
265
266 extern PetscErrorCode process_statistics_add_pdata(process_statistics *,
    process_data *);
267
268 extern PetscErrorCode process_statistics_add_accept(process_statistics *,
    tcpaccept_entry *);
269
270
271 extern PetscErrorCode process_statistics_add_connect(process_statistics *,
    tcpconnect_entry *);
272
273
274 extern PetscErrorCode process_statistics_add_connlat(process_statistics *,
    tcpconnlat_entry *);
275
276
277 extern PetscErrorCode process_statistics_add_life(process_statistics *,
    tcplife_entry *);
278
279
280 extern PetscErrorCode process_statistics_add_retrans(process_statistics *,
    tcpretrans_entry *);
281
282
283 extern PetscErrorCode process_statistics_get_pid_data(process_statistics *,
    PetscInt,
    process_data *);
284
285
286
287
288 typedef struct {
289     FILE *file;
290     long offset;
291 } file_wrapper;
292
293
294 extern long get_file_end_offset(file_wrapper *);
295
296 /* returns 0 if no new data; if the return value is non-zero, it is the difference
    between the new and old end-of-file. However, the file is moved back to the old
    end, so you can read the new data. */
297 extern long has_new_data(file_wrapper *);

```

```
298  
299  
300  
301 #endif
```